


Triton kernel tuned for H=256 LOSES to PyTorch reference at H=384 B=1024 (M10 audit)



Pathfinder-Wide-Long (H=384), d=7, p=0.007 on NVIDIA H200 SXM. Triton kernel block-tile sizes were chosen for H=256 (1.78x speedup at B=1024) and are inefficient at H=384.